

Chathura Dharmasiri

AI/Software Engineer - Intern

itsmechathura@outlook.com | +94 71 748 0048 | Wategama, Kandy, LK
linkedin.com/in/dharmasiri17 | github.com/itschathura | chathuradev.pages.dev

SUMMARY

Development-focused undergraduate pursuing a BSc (Hons) in Computer Science, with hands-on experience building and deploying full-stack and AI-powered applications using Python, Next.js, FastAPI and Docker. Skilled in REST API development, database design and LLM-based systems such as RAG pipelines. Strong problem-solver who learns new technologies quickly and works well in a team. Eager to contribute to a software development team and gain industry experience through a focused internship.

TECHNICAL SKILLS

Languages: Python, TypeScript, JavaScript, SQL, MATLAB

AI & LLMs: RAG, LangChain, LangGraph, Google Gemini API, Groq, Prompt Engineering, Vector Embeddings, Semantic Search, Guardrails, LLM Evaluation

ML & Data: XGBoost, Scikit-learn, Pandas, NumPy, Feature Engineering, Time-Series Analysis, FastF1 Telemetry

Vector & Databases: ChromaDB, PostgreSQL, ScyllaDB (Cassandra CQL)

Backend & Frontend: FastAPI, REST APIs, SSE Streaming, Next.js, React, Tailwind CSS

Tools: Docker, Git, GitHub Actions, Postman, Cloudflare Pages, Vercel, Linux/Ubuntu

PROJECTS & RESEARCH

DocuMind AI - RAG Document Intelligence Platform

Aug 2026

Next.js, FastAPI, Gemini API, ChromaDB, PyMuPDF, PostgreSQL, Docker

[GitHub](#) | [Live Demo](#)

- Built an end-to-end **RAG** platform with **FastAPI** and **Next.js** where users upload multi-page PDFs and get grounded answers and automated summaries.
- Engineered an ingestion pipeline using PyMuPDF with OCR fallback, sliding-window chunking (500 tokens, 10% overlap) and Gemini embeddings, indexed into **ChromaDB** with per-document metadata filtering.
- Reduced hallucinations with constrained context-grounding prompts and source-snippet attribution so every answer links back to the retrieved passages.
- Hardened API security by removing IDOR risks with UUIDv4 document IDs, enforcing MIME/magic-byte file validation (25 MB limit) and isolating storage per client.
- Delivered a responsive UI with live Markdown rendering, source inspection and dark/light mode, deployed on **Cloudflare Pages** and **Docker**.

Inferring Battery State-of-Charge (SoC) and Predicting Overtaking Probability in Formula 1 Racing Using Machine Learning Under 2026 Power Unit Regulations

Aug 2026 - Ongoing

Final-Year Research Project (CSH4216) | Python, XGBoost, ScyllaDB, FastAPI, React 19, FastF1, Docker

[GitHub](#)

- Architected a real-time telemetry streaming and predictive analytics platform that ingests live F1 SignalR timing feeds and predicts overtake probability for the full 20-car grid.
- Trained and evaluated an **XGBoost** classifier on 10M+ rows of multi-round Grand Prix telemetry, handling severe class imbalance with `scale_pos_weight` and PR-AUC across 10 features (closing delta, speed, throttle, braking, DRS, RPM, tyre degradation, energy state).
- Modeled a physics-based battery State-of-Charge engine aligned with the FIA 2026 power unit rules (350 kW MGU-K), simulating four energy-harvesting modes to estimate battery advantage in wheel-to-wheel battles.
- Built a low-latency streaming pipeline on **ScyllaDB** and **AsyncIO** with lockless atomic state swaps (`os.replace`) for high-frequency ingestion without buffer drops.
- Developed a 10 Hz live telemetry dashboard in **React 19** and Tailwind CSS with head-to-head battle gauges, closing-speed and energy-differential meters.

SmartShop AI - Multimodal Shopping Agent

Ongoing

Next.js, FastAPI, LangChain, LangGraph, Gemini, PostgreSQL, Docker, Langfuse, GitHub Actions

[GitHub](#)

- Architected an end-to-end multimodal AI shopping assistant using **FastAPI**, **Next.js** and **LangGraph**, with real-time SSE streaming recommendations and image-based receipt/product search.
- Engineered a 3-tier, zero-cost guardrail pipeline (regex filters plus structured Gemini classifier) to block off-topic and prompt-injection queries before they reach the main LLM.
- Designed a database validation layer on **PostgreSQL (asyncpg)** that checks every recommended product's price and stock to prevent hallucinated items.
- Built multi-turn agentic workflows with LangChain and LangGraph that persist dietary preferences, order history and budget constraints as personalized agent memory.
- Set up CI/CD and evaluation pipelines with **GitHub Actions** and **Langfuse**, benchmarking tool-calling accuracy on golden datasets under free-tier rate limits.

Xtreeme Computer Parts Shop - MERN Stack E-Commerce Web App

2024

MongoDB, Express.js, React, Node.js, REST API, Postman

Personal Project

- Built a full-stack **MERN** web application for browsing and selling computer parts such as processors, graphics cards, keyboards, mice and monitors.
- Developed a RESTful API with **Express.js** and **Node.js** offering full CRUD operations on products (create, read, update and delete by ID), backed by **MongoDB**.
- Designed MongoDB product schemas with category, price, description, image and stock fields.
- Tested and debugged all API endpoints with **Postman** during development.

EDUCATION

University of Vavuniya

BSc (Hons) in Computer Science

Sri Lanka

Sept. 2023 – Present

CERTIFICATIONS

CS50's Introduction to Artificial Intelligence with Python

Harvard University

Sep 2026

[Credential](#)

MATLAB Onramp

MathWorks

Sep 2026

[Credential](#)

Python Programming - Trainee Full Stack Developer

University of Moratuwa

May 2024

Credential ID: WgGERcT8FF | [Verify](#)

SOFT SKILLS

Problem Solving | Teamwork | Adaptability | Time Management | Quick Learner | Can-Do Attitude

REFERENCES

Dr. S. Kirushanth

Senior Lecturer Gr.I, University of Vavuniya

SivaramalingamK@vau.ac.lk

+94 77 672 7201

Anushka Muthusinghe

Associate Technical Lead

anushka.muthusinghe@gmail.com

+94 76 676 8599